

Homework 2

May 1, 2019

1 Homework 2


Data Mining Ross Lewis

Question 1

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np

numImages = 16
fig = plt.figure(figsize=(7,7))
imgData = np.zeros(shape=(numImages,36963))

for i in range(1,numImages+1):
    filename = 'pics/Picture'+str(i)+'.jpg'
    img = mpimg.imread(filename)
    ax = fig.add_subplot(4,4,i)
    plt.imshow(img)
    plt.axis('off')
    ax.set_title(str(i))
    imgData[i-1] = np.array(img.flatten()).reshape(1,img.shape[0]*img.shape[1]*img.shap
```



```
In [2]: import pandas as pd
        from sklearn.decomposition import PCA

        numComponents = 2

        pca = PCA(numComponents)
        pca.fit(imgData)
        projected = pca.transform(imgData)

        projected = pd.DataFrame(projected, columns=['pc1', 'pc2'], index=range(1, numImages+1))
        projected['food'] = ['burger', 'burger', 'burger', 'burger', 'drink', 'drink', 'drink', 'drink',
                             'pasta', 'pasta', 'pasta', 'pasta', 'chicken', 'chicken', 'chicken']
        projected
```

```
Out [2]:
```

	pc1	pc2	food
1	-1576.667016	6641.462674	burger
2	-493.825788	6396.184274	burger
3	990.057233	7236.840769	burger
4	2189.898458	9051.034582	burger
5	-7843.064019	-1061.583683	drink
6	-8498.435161	-5438.442879	drink
7	-11181.787910	-5319.885107	drink
8	-6851.941725	1124.542067	drink
9	7635.140046	-5043.628374	pasta
10	-708.061138	-528.765053	pasta
11	7236.249736	-5301.766423	pasta
12	4417.335795	-4659.524969	pasta
13	11864.507522	1472.188629	chicken
14	76.468838	1366.295780	chicken
15	-7505.666465	-1163.706654	chicken
16	10249.791595	-4771.245634	chicken


Finally, we draw a scatter plot to display the projected values. Observe that the images of burgers, drinks, and pastas are all projected to the same region. However, the images for fried chicken (shown as black squares in the diagram) are harder to discriminate.

Code:

```
In [23]: import matplotlib.pyplot as plt

colors = {'burger':'b', 'drink':'r', 'pasta':'g', 'chicken':'k'}
markerTypes = {'burger': '+', 'drink': 'x', 'pasta': 'o', 'chicken': 's'}

for foodType in markerTypes:
    d = projected[projected['food']==foodType]
    plt.scatter(d['pc1'],d['pc2'],marker=markerTypes[foodType],c=colors[foodType])
```



1.0.1 Extra Credit

In [107]: `import pandas as pd`

```
prof = pd.read_csv('country_profile_variables.csv')
#prof.select_dtypes(['number']).head()
#prof = pd.concat([profDat['Region'],prof.select_dtypes(['number'])],axis=1)
prof.head()
```

```
Out[107]:
```

	country	Region	Surface area (km2)	\
0	Afghanistan	SouthernAsia	652864	
1	Albania	SouthernEurope	28748	
2	Algeria	NorthernAfrica	2381741	
3	American Samoa	Polynesia	199	
4	Andorra	SouthernEurope	468	

	Population in thousands (2017)	Population density (per km2, 2017)	\
0	35530	54.4	
1	2930	106.9	
2	41318	17.3	
3	56	278.2	
4	77	163.8	

	Sex ratio (m per 100 f, 2017)	\
0	106.3	
1	101.9	

2	102.0
3	103.6
4	102.3

	GDP: Gross domestic product (million current US\$)	\
0	20270	
1	11541	
2	164779	
3	-99	
4	2812	

	GDP growth rate (annual %, const. 2005 prices)	\
0	-2.4	
1	2.6	
2	3.8	
3	-99	
4	0.8	

	GDP per capita (current US\$)	Economy: Agriculture (% of GVA)	...	\
0	623.2	23.3	...	
1	3984.2	22.4	...	
2	4154.1	12.2	...	
3	-99.0	-99	...	
4	39896.4	0.5	...	

	Mobile-cellular subscriptions (per 100 inhabitants).1	\
0	8.3	
1	63.3	
2	38.2	
3	-99	
4	96.9	

	Individuals using the Internet (per 100 inhabitants)	\
0	42	
1	130	
2	135	
3	92	
4	13	

	Threatened species (number)	Forested area (% of land area)	\
0	2.1	9.8/0.3	
1	28.2	5.7/2.0	
2	0.8	145.4/3.7	
3	87.9	-99	
4	34.0	0.5/6.4	

	CO2 emission estimates (million tons/tons per capita)	\
0	63	

1	84
2	5900
3	-99
4	1

	Energy production, primary (Petajoules) \
0	5
1	36
2	55
3	-99
4	119

	Energy supply per capita (Gigajoules) \
0	78.2/47.0
1	94.9/95.2
2	84.3/81.8
3	100.0/100.0
4	100.0/100.0

	Pop. using improved drinking water (urban/rural, %) \
0	45.1/27.0
1	95.5/90.2
2	89.8/82.2
3	62.5/62.5
4	100.0/100.0

	Pop. using improved sanitation facilities (urban/rural, %) \
0	21.43
1	2.96
2	0.05
3	-99
4	-99

	Net Official Development Assist. received (% of GNI)
0	-99
1	-99
2	-99
3	-99
4	-99

[5 rows x 50 columns]

```
In [108]: from sklearn.preprocessing import normalize
```

```
numComponents = 2
```

```
pca = PCA(numComponents)
```

```
#data = normalize(prof.select_dtypes(['number']), axis=0, norm='max')
```

```

pca.fit(normalize(prof.select_dtypes(['number']), axis=0, norm='max'))
projected = pca.transform(normalize(prof.select_dtypes(['number']), axis=0, norm='ma

projected = pd.DataFrame(projected,columns=['pc1','pc2'],index=range(0,len(prof)))
projected['region'] = prof['Region']
projected.head()

```

```

Out[108]:
   pc1    pc2    region
0 -1.225279  0.101471  SouthernAsia
1 -1.080330  0.048536  SouthernEurope
2 -1.221848 -0.069134  NorthernAfrica
3  5.664929  0.996393    Polynesia
4 -1.125041  0.017737  SouthernEurope

```

```


In [115]: print(len(np.unique(projected['region'])))

for reg in np.unique(projected['region']):
    d = projected[projected['region']==reg]

    plt.scatter(d['pc1'],d['pc2'])

```

22



Above is the normalized data projected onto two dimensions with different regions labeled by color.

Question 2

1.1 Decision Tree

```
In [24]: # Load the dataset
```

```
import pandas as pd
```

```
data = pd.read_csv('vertebrate.csv', header='infer')  
data
```

```
Out[24]:
```

	Name	Warm-blooded	Gives Birth	Aquatic Creature	\
0	human	1	1	0	
1	python	0	0	0	
2	salmon	0	0	1	
3	whale	1	1	1	
4	frog	0	0	1	
5	komodo	0	0	0	
6	bat	1	1	0	
7	pigeon	1	0	0	
8	cat	1	1	0	
9	leopard shark	0	1	1	
10	turtle	0	0	1	
11	penguin	1	0	1	
12	porcupine	1	1	0	
13	eel	0	0	1	
14	salamander	0	0	1	

	Aerial Creature	Has Legs	Hibernates	Class
0	0	1	0	mammals
1	0	0	1	reptiles
2	0	0	0	fishes
3	0	0	0	mammals
4	0	1	1	amphibians
5	0	1	0	reptiles
6	1	1	1	mammals
7	1	1	0	birds
8	0	1	0	mammals
9	0	0	0	fishes
10	0	1	0	reptiles
11	0	1	0	birds
12	0	1	1	mammals
13	0	0	0	fishes
14	0	1	1	amphibians

```
In [28]: # Pre-processing data
```

```
# convert the classes to be binary: replacing fishes, birds, amphibians and reptiles a
```



```
data['Class'] = np.where(data['Class'] == 'mammals', 'mammals', 'non-mammals')
data
```

```
Out[28]:
```

	Name	Warm-blooded	Gives Birth	Aquatic Creature	\
0	human	1	1	0	
1	python	0	0	0	
2	salmon	0	0	1	
3	whale	1	1	1	
4	frog	0	0	1	
5	komodo	0	0	0	
6	bat	1	1	0	
7	pigeon	1	0	0	
8	cat	1	1	0	
9	leopard shark	0	1	1	
10	turtle	0	0	1	
11	penguin	1	0	1	
12	porcupine	1	1	0	
13	eel	0	0	1	
14	salamander	0	0	1	

	Aerial Creature	Has Legs	Hibernates	Class
0	0	1	0	mammals
1	0	0	1	non-mammals
2	0	0	0	non-mammals
3	0	0	0	mammals
4	0	1	1	non-mammals
5	0	1	0	non-mammals
6	1	1	1	mammals
7	1	1	0	non-mammals
8	0	1	0	mammals
9	0	0	0	non-mammals
10	0	1	0	non-mammals
11	0	1	0	non-mammals
12	0	1	1	mammals
13	0	0	0	non-mammals
14	0	1	1	non-mammals

```
In [29]: # Pandas provides a function cross-tabulation that can examine the relationship between
pd.crosstab([data['Warm-blooded'], data['Gives Birth']], data['Class'])
```

```
Out[29]:
```

Class		mammals	non-mammals
Warm-blooded	Gives Birth		
0	0	0	7
	1	0	1
1	0	0	2
	1	5	0

The results above show that it is possible to distinguish mammals from non-mammals using

these two attributes alone since each combination of their attribute values would yield only instances that belong to the same class. For example, mammals can be identified as warm-blooded vertebrates that give birth to their young. Such a relationship can also be derived using a decision tree classifier, as shown by the example given in the next subsection.

```
In [58]: # apply a decision tree classifier to the vertebrate dataset described in the previous
```

```
from sklearn import tree

Y = data['Class']
X = data.drop(['Name', 'Class'], axis=1)

clf = tree.DecisionTreeClassifier(criterion='gini', splitter='random', max_depth=3, min_samples_split=2)
clf = clf.fit(X, Y)
```

The preceding commands will extract the predictor (X) and target class (Y) attributes from the vertebrate dataset and create a decision tree classifier object using entropy as its impurity measure for splitting criterion. The decision tree class in Python sklearn library also supports using 'gini' as impurity measure. The classifier above is also constrained to generate trees with a maximum depth equals to 3. Next, the classifier is trained on the labeled data using the fit() function.

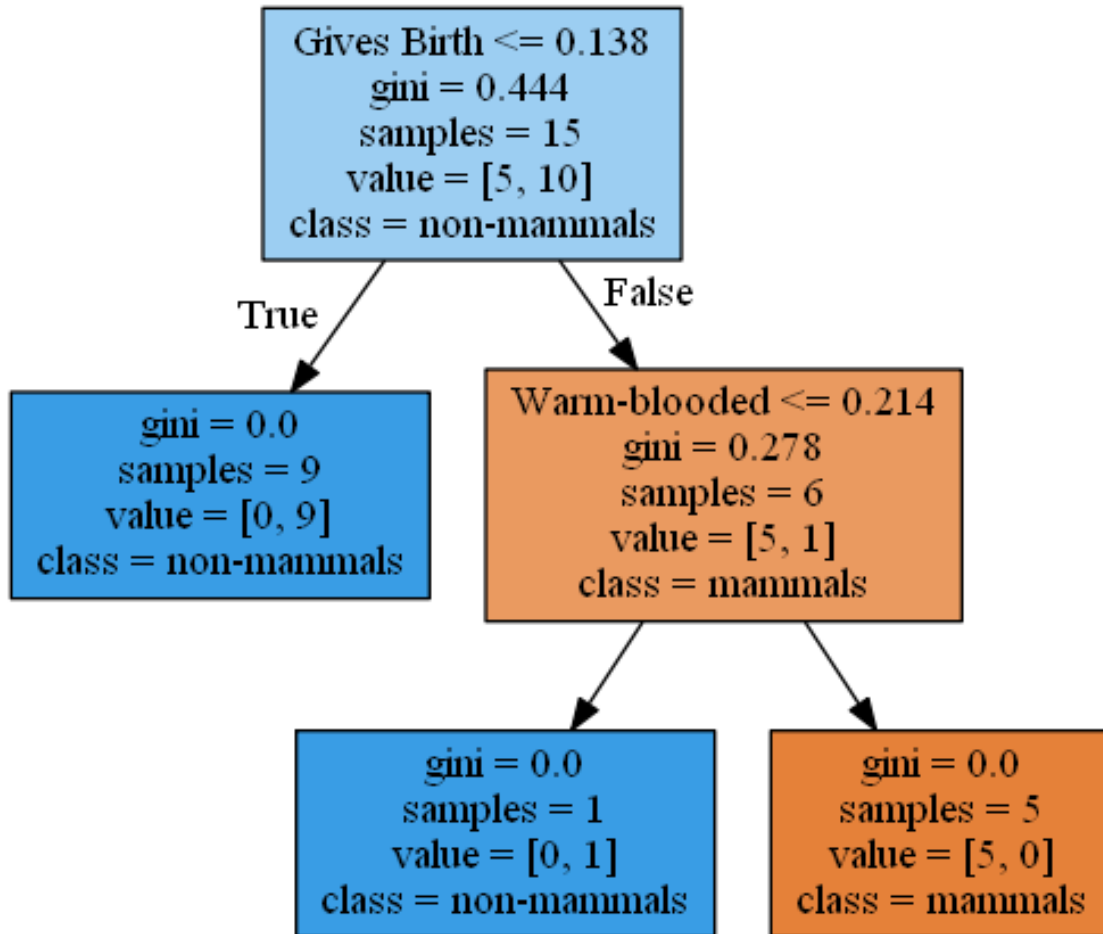
We can plot the resulting decision tree obtained after training the classifier. To do this, you must first install both graphviz (<http://www.graphviz.org>) and its Python interface called pydotplus (<http://pydotplus.readthedocs.io/>).

```
In [59]: # visualizing the tree
```

```
import pydotplus
from IPython.display import Image

dot_data = tree.export_graphviz(clf, feature_names=X.columns, class_names=['mammals', 'non-mammals'],
                                out_file=None)
graph = pydotplus.graph_from_dot_data(dot_data)
Image(graph.create_png())
```

```
Out [59]:
```



Next, suppose we apply the decision tree to classify the following test examples.

```
In [60]: testData = [['gila monster',0,0,0,0,1,1,'non-mammals'],
                    ['platypus',1,0,0,0,1,1,'mammals'],
                    ['owl',1,0,0,1,1,0,'non-mammals'],
                    ['dolphin',1,1,1,0,0,0,'mammals']]
testData = pd.DataFrame(testData, columns=data.columns)
testData
```

```
Out [60]:
```

	Name	Warm-blooded	Gives Birth	Aquatic Creature	Aerial Creature	\
0	gila monster	0	0	0	0	
1	platypus	1	0	0	0	
2	owl	1	0	0	1	
3	dolphin	1	1	1	0	

	Has Legs	Hibernates	Class
0	1	1	non-mammals
1	1	1	mammals
2	1	0	non-mammals
3	0	0	mammals

We first extract the predictor and target class attributes from the test data and then apply the decision tree classifier to predict their classes.

```
In [62]: testY = testData['Class']
         testX = testData.drop(['Name', 'Class'], axis=1)

         predY = clf.predict(testX)
         predictions = pd.concat([testData['Name'], pd.Series(predY, name='Predicted Class')], axis=1)
         predictions
```

```
Out[62]:
```

	Name	Predicted Class
0	gila monster	non-mammals
1	platypus	non-mammals
2	owl	non-mammals
3	dolphin	mammals

Except for platypus, which is an egg-laying mammal, the classifier correctly predicts the class label of the test examples. We can calculate the accuracy of the classifier on the test data as shown by the example given below.

```
In [63]: from sklearn.metrics import accuracy_score

         print('Accuracy on test data is %.2f' % (accuracy_score(testY, predY)))
```

```
Accuracy on test data is 0.75
```