

Ross Lewis

## 425 - Machine Learning – Homework 4

### PART 1

#### Problem I:

I tried three models. The first has 1 hidden layer, uses relu as an activation function, and has a softmax output. The second has 2 hidden layers and uses sigmoid activation and output functions. The last has 3 hidden layers, uses relu as an activation function, and softmax as an output function.

```
#Model 1
#1 hidden Layer
#relu activation
#softmax output
model1 = Sequential()
model1.add(Dense(128, activation='relu', input_shape=(784,)))
model1.add(Dense(num_classes, activation='softmax'))

model1.compile(loss='categorical_crossentropy',
               optimizer=RMSprop(),
               metrics=['accuracy'])
print(model1.summary())

model1.fit(xtrain, ytrain,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1)

ypred = model1.predict_classes(xval)

conf, acc, rec, prec = func_confusion_matrix(yval, ypred)
print(conf)
print(acc)
print(rec)
print(prec)
```

```
[[ 951  0  1  2  1  0  0  2  3  0]
 [ 0 1035  5  0  1  0  0  3  8  2]
 [ 1  2 987  6  1  0  2  7  3  1]
 [ 1  0  9 1001  2  9  0  5  3  2]
 [ 2  0  3  0 1017  0  4  4  1 11]
 [ 0  1  1  9  2 852  3  1  6  2]
 [ 7  1  2  0  2  3 953  0  6  0]
 [ 1  3  4  1  2  0  0 1035  1  6]
 [ 4  2  4  7  0  3  0  3 992  3]
```

[ 1 1 1 6 7 3 0 9 5 947]]

Model 1 accuracy = 0.977

```
#model 2
#2 hidden layers
#sigmoid activation
#sigmoid output
model2 = Sequential()
model2.add(Dense(128, activation='sigmoid', input_shape=(784,)))
model2.add(Dropout(0.2))
model2.add(Dense(64, activation='sigmoid'))
model2.add(Dense(num_classes, activation='sigmoid'))

model2.compile(loss='categorical_crossentropy',
               optimizer=RMSprop(),
               metrics=['accuracy'])

model2.fit(xtrain, ytrain,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1)

ypred = model2.predict_classes(xval)

conf, acc, rec, prec = func_confusion_matrix(yval, ypred)
print(conf)
print(acc)
print(rec)
print(prec)
```

[[ 944 0 1 3 1 2 3 1 5 0]

[ 0 1033 5 1 2 0 0 2 6 5]

[ 2 1 976 6 5 0 5 4 10 1]

[ 2 0 12 975 1 24 0 10 5 3]

[ 3 0 3 0 1007 0 6 3 1 19]

[ 2 3 2 14 3 838 4 2 5 4]

[ 4 2 2 1 2 5 954 0 4 0]

[ 2 6 2 3 1 0 0 1032 0 7]

[ 1 4 7 10 0 5 2 2 983 4]

[ 2 2 1 8 8 4 0 12 8 935]]

Model 2 accuracy = 0.9677

```

#model 3
#3 hidden layers
#relu activation
#softmax output
model3 = Sequential()
model3.add(Dense(256, activation='relu', input_shape=(784,)))
model3.add(Dropout(0.2))
model3.add(Dense(128, activation='relu'))
model3.add(Dropout(0.2))
model3.add(Dense(64, activation='relu'))
model3.add(Dense(num_classes, activation='softmax'))

model3.compile(loss='categorical_crossentropy',
               optimizer=RMSprop(),
               metrics=['accuracy'])

model3.fit(xtrain, ytrain,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1)

ypred = model3.predict_classes(xval)

conf, acc, rec, prec = func_confusion_matrix(yval, ypred)
print(conf)
print(acc)
print(rec)
print(prec)

[[ 955  0  2  1  0  0  1  0  1  0]
 [ 0 1041  3  0  2  0  1  1  3  3]
 [ 2  1 988  8  1  0  2  2  5  1]
 [ 0  0  7 1012  0  7  0  2  2  2]
 [ 1  0  3  0 1016  0  2  1  1 18]
 [ 1  1  1  9  1 857  2  0  3  2]
 [ 4  1  1  0  1  3 960  0  4  0]
 [ 0  4  0  3  2  0  0 1030  1 13]
 [ 1  1  4 10  0  3  0  2 992  5]
 [ 1  1  0  3  6  6  0  4  2 957]]

```

Model 3 accuracy = 0.9808

Model 3 has the best accuracy.

Question II:

Confusion Matrix for model 3 on the test data:

```
[[ 971  1  1  0  1  1  2  1  2  0]
```

```
[ 0 1125 3 1 0 0 2 0 4 0]
[ 0 0 1016 4 2 0 2 4 4 0]
[ 0 0 4 995 0 2 0 4 3 2]
[ 2 0 1 1 958 0 3 1 1 15]
[ 2 0 0 13 1 869 3 0 2 2]
[ 3 2 0 0 3 7 942 0 1 0]
[ 2 6 9 6 1 0 0 988 3 13]
[ 1 1 3 6 1 4 0 2 953 3]
[ 3 2 0 8 4 1 0 1 1 989]]
```

Accuracy of model 3:


0.9806

Recall of model 3:


```
[0.99081633 0.99118943 0.98449612 0.98514851 0.97556008 0.97421525
0.98329854 0.96108949 0.97843943 0.98017839]
```

Precision of model 3:

```
[0.98678862 0.98944591 0.97974928 0.9622824 0.98661174 0.98303167
0.98742138 0.98701299 0.97843943 0.96582031]
```




predicted 8  
actual 3




predicted 9  
actual 4

If the lines were connected, the three would look a lot like an 8. Same with the top of the 4 and being similar to a 9.




predicted 2  
actual 4




predicted 7  
actual 2

The 4 is quite ambiguous so it would be hard to predict. This 2 looks similar to a 7.




predicted 3  
actual 5




predicted 7  
actual 3

This 5 looks like a 3 without the top part, and this 3 looks like a 7.




predicted 0  
actual 6




predicted 9  
actual 4

The 6 is almost a 0, and the 4 could be a 9 or a 4.





predicted 2  
actual 8




predicted 2  
actual 8


This 8 seems like a bad prediction, but the bottom 8 isn't even drawn well.

## PART 2

A C value of 5 looks like it is the best for this data.



Linear performs better than poly and rbf.



Metrics for the test set:

**Confusion Matrix:**

```
[[52  0]
 [ 8 40]]
```

**Average Accuracy: 0.92**

**Per-Class Precision: [0.86666667 1. ]**

**Per-Class Recall: [1. 0.83333333]**

Here are 5 incorrect predictions and 5 correct predictions. Correct predictions tend to have more extreme value features where the incorrect predictions are in the middle.

predicted 1.0

actual -1.0

data [ 1. 11.1 9.9 23.8 27.1 9.8]

predicted 1.0

actual -1.0

data [ 1. 12.3 11. 26.8 31.5 11.4]

predicted 1.0

actual -1.0

data [ 1. 9.2 7.8 19. 22.4 7.7]

predicted 1.0

actual -1.0

data [ 0. 9.1 6.9 16.7 18.6 7.4]

predicted 1.0

actual -1.0

data [ 1. 12.8 10.9 27.4 31.5 11. ]

predicted -1.0

actual -1.0

data [ 1. 13.9 11.1 29.2 33.3 12.1]

predicted -1.0

actual -1.0

data [ 1. 19.8 14.2 43.2 49.7 18.6]

predicted -1.0

actual -1.0

data [ 1. 19.7 15.3 41.9 48.5 17.8]

predicted 1.0

actual 1.0

data [ 1. 14.7 12.5 30.1 34.7 12.5]

predicted 1.0

actual 1.0

data [ 0. 15.7 13.6 31. 34.8 13.8]